

roboflow /notebooks

How to Train YOLOv8 Object Detection on a Custom Dataset

 Roboflow Blog  Youtube  GitHub

Ultralytics YOLOv8 is the latest version of the YOLO (You Only Look Once) object detection and image segmentation model developed by Ultralytics. The YOLOv8 model is designed to be fast, accurate, and easy to use, making it an excellent choice for a wide range of object detection and image segmentation tasks. It can be trained on large datasets and is capable of running on a variety of hardware platforms, from CPUs to GPUs.

Disclaimer

YOLOv8 is still under heavy development. Breaking changes are being introduced almost weekly. We strive to make our YOLOv8 notebooks work with the latest version of the library. Last tests took place on **03.01.2024** with version **YOLOv8.0.196**.

If you notice that our notebook behaves incorrectly - especially if you experience errors that prevent you from going through the tutorial - don't hesitate! Let us know and open an [issue](#) on the Roboflow Notebooks repository.

Accompanying Blog Post

We recommend that you follow along in this notebook while reading the blog post on how to train YOLOv8 Object Detection, concurrently.

Pro Tip: Use GPU Acceleration

If you are running this notebook in Google Colab, navigate to `Edit -> Notebook settings -> Hardware accelerator`, set it to `GPU`, and then click `Save`. This will ensure your notebook uses a GPU, which will significantly speed up model training times.

Steps in this Tutorial

In this tutorial, we are going to cover:

- Before you start
- Install YOLOv8
- CLI Basics
- Inference with Pre-trained COCO Model
- Roboflow Universe
- Preparing a custom dataset
- Custom Training
- Validate Custom Model
- Inference with Custom Model

Let's begin!

▼ Before you start

Let's make sure that we have access to GPU. We can use `nvidia-smi` command to do that. In case of any problems navigate to `Edit -> Notebook settings -> Hardware accelerator`, set it to `GPU`, and then click `Save`.

```
!nvidia-smi
```

```
Wed May 15 19:46:51 2024
```

| NVIDIA-SMI 535.104.05 | | | Driver Version: 535.104.05 | | CUDA Version | |
|-----------------------|----------|---------------|----------------------------|-----------------|--------------|--|
| GPU | Name | Persistence-M | Bus-Id | Disp.A | Volatile U | |
| Fan | Temp | Pwr:Usage/Cap | | Memory-Usage | GPU-Util | |
| 0 | Tesla T4 | Off | 00000000:00:04.0 | Off | | |
| N/A | 71C P8 | 11W / 70W | | 0MiB / 15360MiB | 0% | |

| Processes: | | | | | | |
|----------------------------|----|----|-----|------|--------------|--|
| GPU | GI | CI | PID | Type | Process name | |
| ID | | ID | | | | |
| No running processes found | | | | | | |

```
import os  
HOME = os.getenv('HOME')
```

```
HOME = os.getenv()  
print(HOME)  
/content
```

▼ Install YOLOv8

⚠️ YOLOv8 is still under heavy development. Breaking changes are being introduced almost weekly. We strive to make our YOLOv8 notebooks work with the latest version of the library. Last tests took place on **03.01.2024** with version **YOLOv8.0.196**.

If you notice that our notebook behaves incorrectly - especially if you experience errors that prevent you from going through the tutorial - don't hesitate! Let us know and open an [issue](#) on the Roboflow Notebooks repository.

YOLOv8 can be installed in two ways-from the source and via pip. This is because it is the first iteration of YOLO to have an official package.

```
# Pip install method (recommended)  
  
!pip install ultralytics==8.0.196  
  
from IPython import display  
display.clear_output()  
  
import ultralytics  
ultralytics.checks()  
  
Ultralytics YOLOv8.0.196 🚀 Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4  
Setup complete ✅ (2 CPUs, 12.7 GB RAM, 30.8/78.2 GB disk)  
  
# Git clone method (for development)  
  
# %cd {HOME}  
# !git clone github.com/ultralytics/ultralytics  
# %cd {HOME}/ultralytics  
# !pip install -e .  
  
# from IPython import display  
# display.clear_output()  
  
# import ultralytics  
# ultralytics.checks()  
  
from ultralytics import YOLO  
  
from IPython.display import display, Image
```

▼ CLI Basics

If you want to train, validate or run inference on models and don't need to make any modifications to the code, using YOLO command line interface is the easiest way to get started. Read more about CLI in [Ultralytics YOLO Docs](#).

```
yolo task=detect    mode=train    model=yolov8n.yaml    args...
      classify      predict      yolov8n-cls.yaml  args...
      segment        val        yolov8n-seg.yaml  args...
                  export      yolov8n.pt       format=onnx  args...
```

▼ Inference with Pre-trained COCO Model

▼ CLI

`yolo mode=predict` runs YOLOv8 inference on a variety of sources, downloading models automatically from the latest YOLOv8 release, and saving results to `runs/predict`.

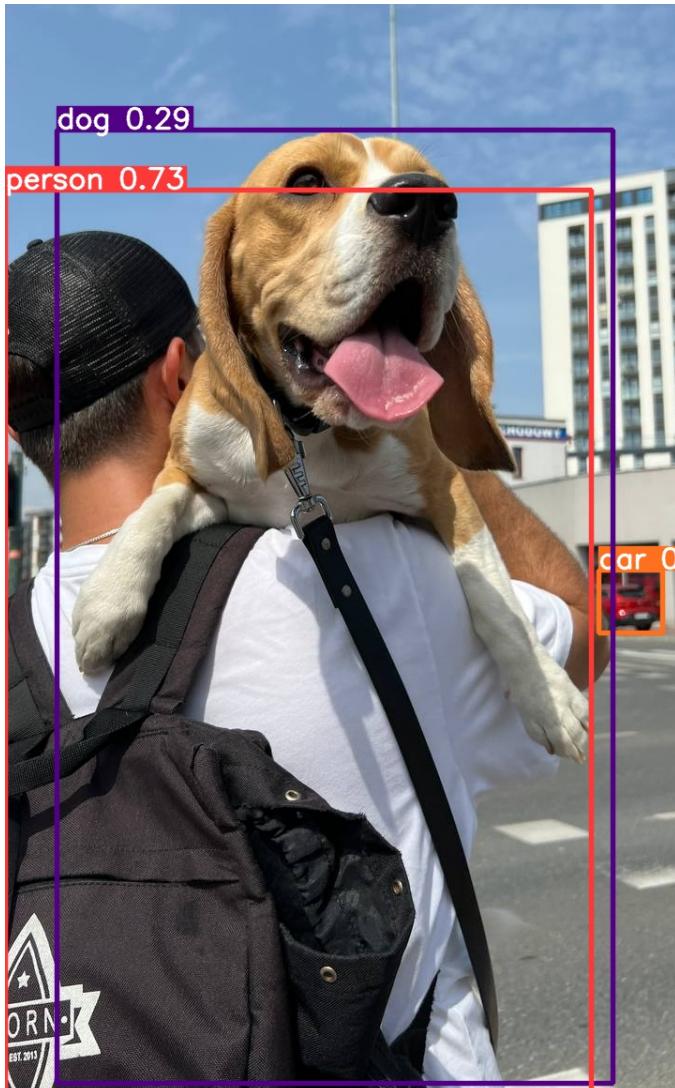
```
%cd {HOME}
!yolo task=detect mode=predict model=yolov8n.pt conf=0.25 source='https://media.r
 /content
Downloading https://github.com/ultralytics/assets/releases/download/v0.0.0/yolov8n.pt
100% 6.23M/6.23M [00:00<00:00, 86.1MB/s]
Ultralytics YOLOv8.0.196 🚀 Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4
YOLOv8n summary (fused): 168 layers, 3151904 parameters, 0 gradients, 8.7 GFL

Found https://media.roboflow.com/notebooks/examples/dog.jpeg locally at dog.jpeg
WARNING ! NMS time limit 0.550s exceeded
image 1/1 /content/dog.jpeg: 640x384 1 person, 1 car, 1 dog, 213.5ms
Speed: 12.9ms preprocess, 213.5ms inference, 1237.1ms postprocess per image
Results saved to runs/detect/predict4
💡 Learn more at https://docs.ultralytics.com/modes/predict
```

```
%cd {HOME}
Image(filename='runs/detect/predict/dog.jpeg', height=600)
```

```
/content
```





▼ Python SDK

The simplest way of simply using YOLOv8 directly in a Python environment.

```
model = YOLO(f'{HOME}/yolov8n.pt')
results = model.predict(source='https://media.roboflow.com/notebooks/examples/doc

Found https://media.roboflow.com/notebooks/examples/dog.jpeg locally at dog.j
image 1/1 /content/dog.jpeg: 640x384 1 person, 1 car, 1 dog, 143.0ms
Speed: 12.6ms preprocess, 143.0ms inference, 511.1ms postprocess per image at

results[0].boxes.xyxy
tensor([[ 0.0000,  314.4717,  625.0754, 1278.1946],
       [ 55.1731,  250.0220,  648.1080, 1266.2720],
       [ 633.2291,  719.5391,  701.0538,  786.0336]], device='cuda:0')

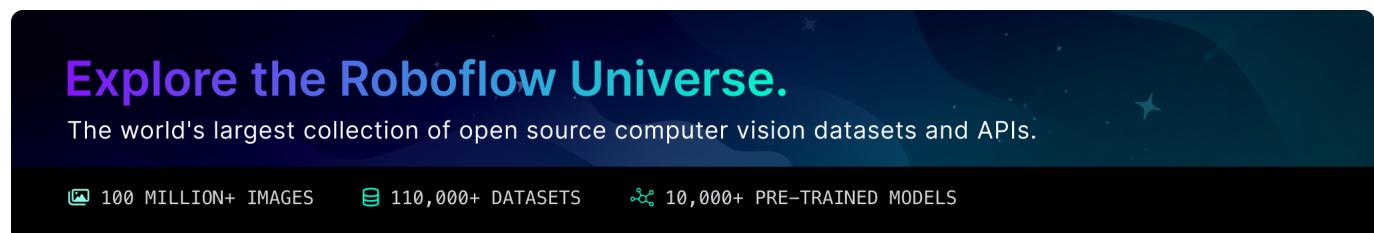
results[0].boxes.conf
```

```
tensor([0.7271, 0.2907, 0.2846], device='cuda:0')
```

```
results[0].boxes.cls  
tensor([ 0., 16., 2.], device='cuda:0')
```

Roboflow Universe

Need data for your project? Before spending time on annotating, check out Roboflow Universe, a repository of more than 110,000 open-source datasets that you can use in your projects. You'll find datasets containing everything from annotated cracks in concrete to plant images with disease annotations.



▼ Preparing a custom dataset

Building a custom dataset can be a painful process. It might take dozens or even hundreds of hours to collect images, label them, and export them in the proper format. Fortunately, Roboflow makes this process as straightforward and fast as possible. Let me show you how!

Step 1: Creating project

Before you start, you need to create a Roboflow [account](#). Once you do that, you can create a new project in the Roboflow [dashboard](#). Keep in mind to choose the right project type. In our case, Object Detection.

The screenshot shows the Roboflow dashboard. At the top, there are navigation links: "Jumpstart", "Projects", "Universe", "Documentation", "Forum", and user info "Piotr Skalski". Below the navigation, there's a "WORKSPACES" section with "Robotflow" (8), "Mohamed Traore" (25), and "Robotflow Universe Proj..." (26). A "Create New Project" button is visible. The main area displays several project cards:

- "creacks": INSTANCE-SEGMENTATION, Modified 7 days ago, 1 CLASSES, 4028 IMAGES
- "test": OBJECT-DETECTION, Modified 21 days ago, 0 TEST, 0 IMAGES
- "football-players-detection": OBJECT-DETECTION, Modified a month ago, 4 CLASSES, 560 IMAGES, 1 MODELS
- "football-players-detection": OBJECT-DETECTION, Modified a month ago, 2 CLASSES, 560 IMAGES
- "doge": INSTANCE-SEGMENTATION, Modified 2 months ago, 1 DOGE, 1 IMAGES

RESOURCES

- Getting Started »
- Tutorials »
- Public Datasets & Models »
- Model Library »
- Help & Support »

Step 2: Uploading images

Next, add the data to your newly created project. You can do it via API or through our [web interface](#).

If you drag and drop a directory with a dataset in a supported format, the Roboflow dashboard will automatically read the images and annotations together.

roboflow Projects Universe Documentation Forum Piotr Skalski

Upload

(?) Want to modify or change classes on your uploaded images?

Batch Name

All Images 0 Annotated 0 Not Annotated 0

ROBOFLOW

Football Players Detection Object Detection

Overview Upload Assign beta Annotate Dataset 0 Generate Versions

Drag and drop images and annotations

Select Files Select Folder

Images jpg, png, bmp Annotations in 26 formats Video mov, mp4, avi

Step 3: Labeling

If you only have images, you can label them in [Roboflow Annotate](#).

← Football Players Detection > Annotate 2 / 560 ...

Annotations

Group: football-players

CLASSES LAYERS

UNUSED CLASSES
goalkeeper
player

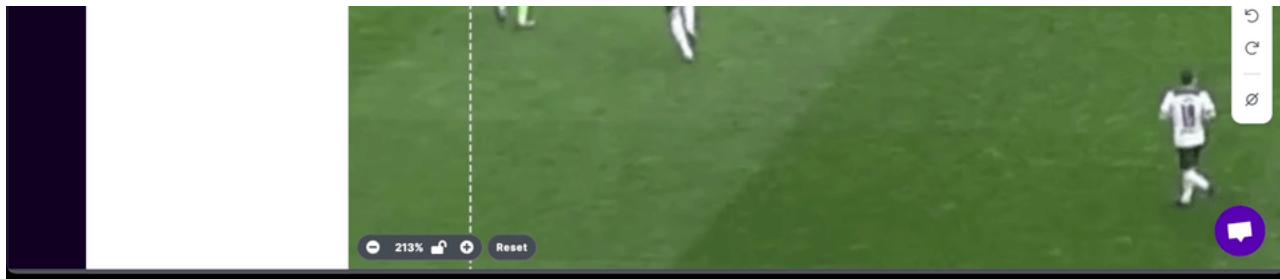
Annotations

ANNOTATION SHORTCUTS

Annotations

Attributes

Raw Data

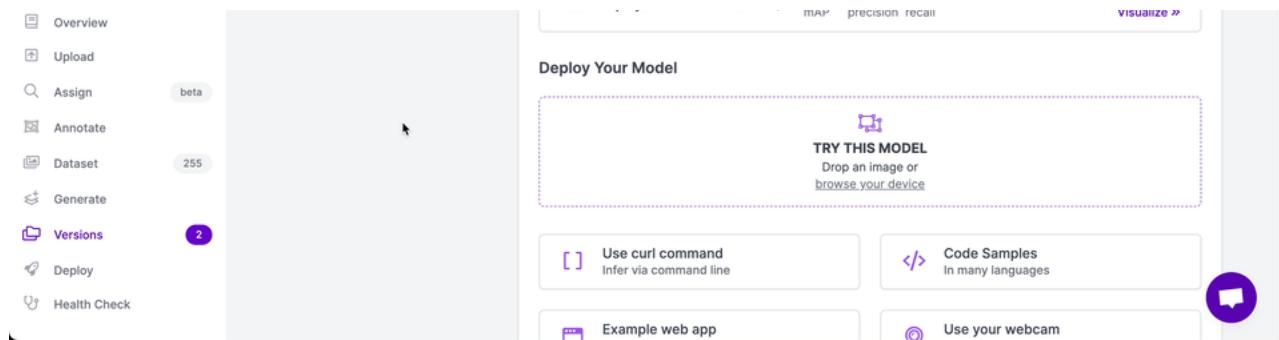


Step 4: Generate new dataset version

Now that we have our images and annotations added, we can Generate a Dataset Version. When Generating a Version, you may elect to add preprocessing and augmentations. This step is completely optional, however, it can allow you to significantly improve the robustness of your model.

Step 5: Exporting dataset

Once the dataset version is generated, we have a hosted dataset we can load directly into our notebook for easy training. Click Export and select the YOLO v5 PyTorch dataset format.



```
!mkdir {HOME}/datasets
%cd {HOME}/datasets

!pip install roboflow --quiet

from roboflow import Roboflow

rf = Roboflow(api_key="6n4vryXJRgiaPTgLHstB")
project = rf.workspace("personal-ubphn").project("klt_detection")
version = project.version(1)
dataset = version.download("yolov8")

mkdir: cannot create directory '/content/datasets': File exists
/content/datasets
loading Roboflow workspace...
loading Roboflow project...
```

Custom Training

```
%cd {HOME}

!yolo task=detect mode=train model=yolov8n.pt data={dataset.location}/data.yaml ep

/content
New https://pypi.org/project/ultralytics/8.2.16 available 😊 Update with 'pip
Ultralytics YOLOv8.0.196 🚀 Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4
engine/trainer: task=detect, mode=train, model=yolov8n.pt, data=/content/data
2024-05-15 19:47:37.882952: E external/local_xla/xla/stream_executor/cuda/cud
2024-05-15 19:47:37.883027: E external/local_xla/xla/stream_executor/cuda/cud
2024-05-15 19:47:37.884920: E external/local_xla/xla/stream_executor/cuda/cud
Overriding model.yaml nc=80 with nc=2
```

| | from | n | params | module |
|---|------|---|--------|----------------------------------|
| 0 | -1 | 1 | 464 | ultralytics.nn.modules.conv.Conv |
| 1 | -1 | 1 | 4672 | ultralytics.nn.modules.conv.Conv |
| 2 | -1 | 1 | 7360 | ultralytics.nn.modules.block.C2f |
| 3 | -1 | 1 | 18560 | ultralytics.nn.modules.conv.Conv |
| 4 | -1 | 2 | 49664 | ultralytics.nn.modules.block.C2f |
| 5 | -1 | 1 | 73984 | ultralytics.nn.modules.conv.Conv |
| 6 | -1 | 2 | 197632 | ultralytics.nn.modules.block.C2f |

```

7          -1  1    295424 ultralytics.nn.modules.conv.Conv
8          -1  1    460288 ultralytics.nn.modules.block.C2f
9          -1  1    164608 ultralytics.nn.modules.block.SPPF
10         -1  1      0 torch.nn.modules.upsampling.Upsample
11        [-1, 6] 1      0 ultralytics.nn.modules.conv.Concat
12         -1  1   148224 ultralytics.nn.modules.block.C2f
13         -1  1      0 torch.nn.modules.upsampling.Upsample
14        [-1, 4] 1      0 ultralytics.nn.modules.conv.Concat
15         -1  1    37248 ultralytics.nn.modules.block.C2f
16         -1  1    36992 ultralytics.nn.modules.conv.Conv
17        [-1, 12] 1      0 ultralytics.nn.modules.conv.Concat
18         -1  1   123648 ultralytics.nn.modules.block.C2f
19         -1  1   147712 ultralytics.nn.modules.conv.Conv
20        [-1, 9] 1      0 ultralytics.nn.modules.conv.Concat
21         -1  1    493056 ultralytics.nn.modules.block.C2f
22       [15, 18, 21] 1   751702 ultralytics.nn.modules.head.Detect
Model summary: 225 layers, 3011238 parameters, 3011222 gradients, 8.2 GFLOPs

```

Transferred 319/355 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs/detect/train4', view at [ht](#)
Freezing layer 'model.22.dfl.conv.weight'

AMP: running Automatic Mixed Precision (AMP) checks with YOL0v8n...

WARNING ! NMS time limit 0.550s exceeded

AMP: checks passed ✓

WARNING ! updating to 'imgsz=640'. 'train' and 'val' imgsz must be an integer

train: Scanning /content/datasets/klt_detection-1/train/labels.cache... 203 i

albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limi

/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork
self.pid = os.fork()

val: Scanning /content/datasets/klt_detection-1/valid/labels.cache... 19 imag
Plotting labels to runs/detect/train4/labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' a
optimizer: AdamW(lr=0.001667, momentum=0.9) with parameter groups 57 weight(d

Image sizes 640 train, 640 val

Using 2 dataloader workers

Logging results to **runs/detect/train4**

Starting training for 100 epochs...

| Epoch | GPU_mem | box_loss | cls_loss | dfl_loss | Instances | Size |
|-------|---------|----------|-----------|----------|-----------|------|
| 1/100 | 2.39G | 1.206 | 3.242 | 1.343 | 21 | 640 |
| | Class | Images | Instances | | Box(P) | R |

!ls {HOME}/runs/detect/train/

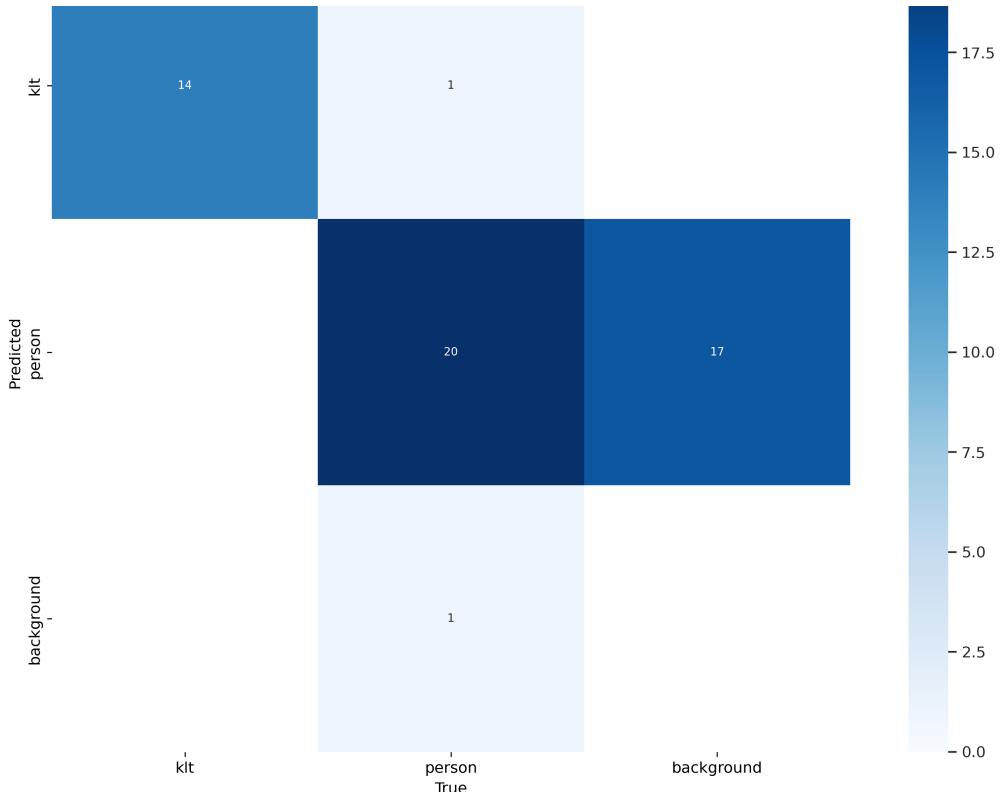
| | | |
|--|------------------|---------|
| args.yaml | labels.jpg | train_b |
| events.out.tfevents.1715801731.4109b3c2dbac.1127.0 | results.csv | train_b |
| labels_correlogram.jpg | train_batch0.jpg | weights |

%cd {HOME}

Image(filename=f'{HOME}/runs/detect/train4/confusion_matrix.png', width=600)

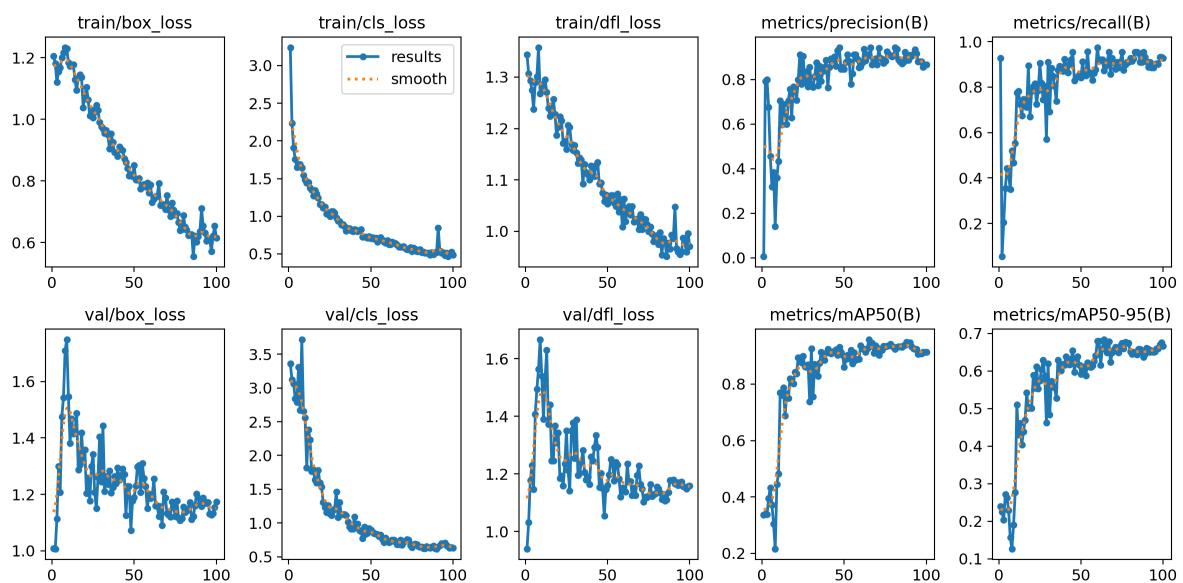
/content





```
%cd {HOME}
Image(filename=f'{HOME}/runs/detect/train4/results.png', width=600)
```

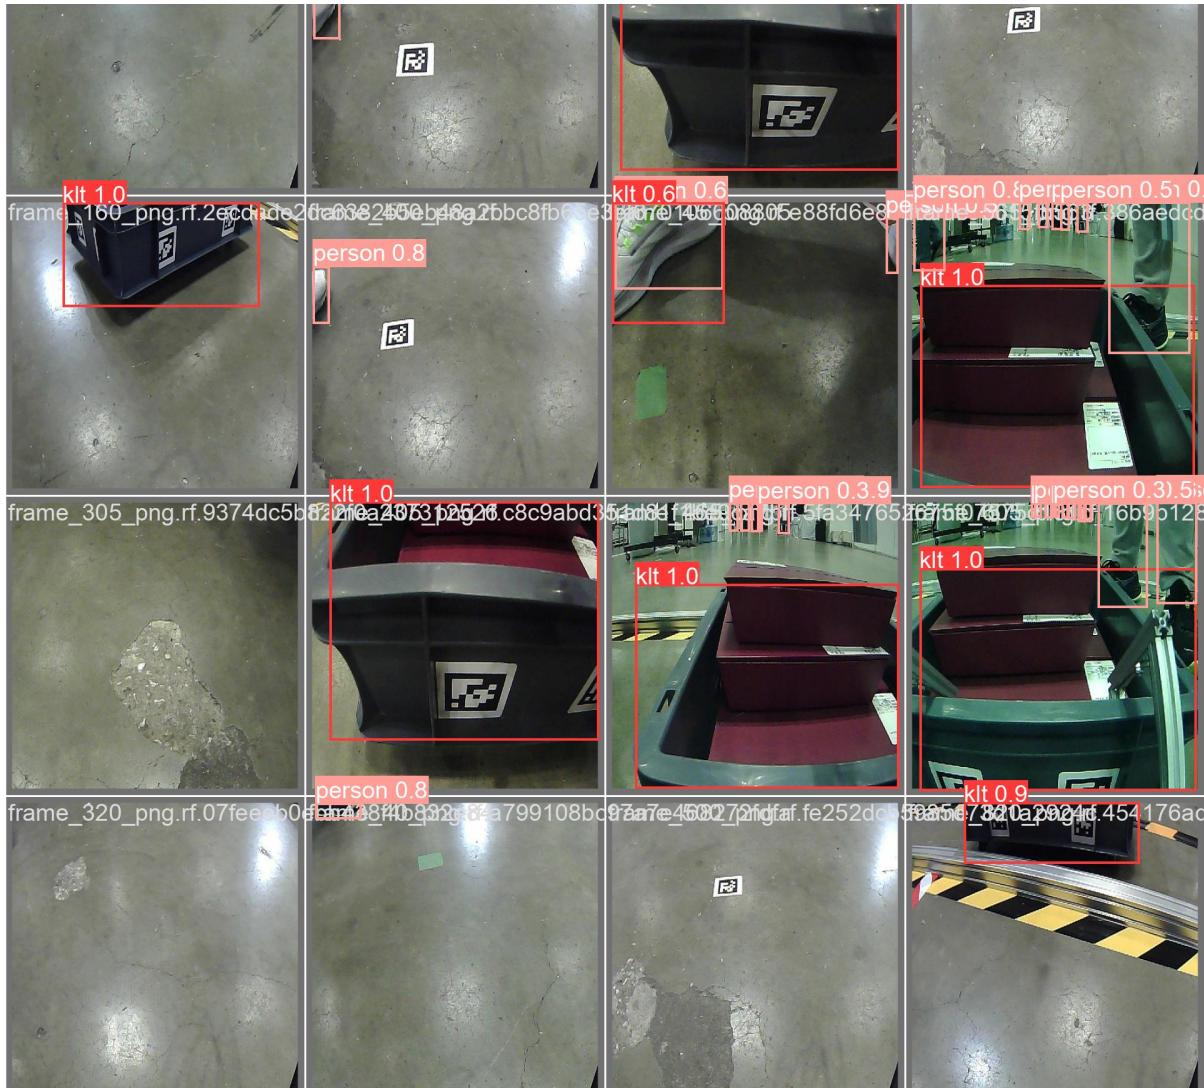
/content



```
%cd {HOME}
Image(filename=f'{HOME}/runs/detect/train4/val_batch0_pred.jpg', width=600)
```

/content





▼ Validate Custom Model

```
%cd {HOME}
```

```
!yolo task=detect mode=val model={HOME}/runs/detect/train4/weights/best.pt data={/content
Ultralytics YOLOv8.0.196 🚀 Python-3.10.12 torch-2.2.1+cu121 CUDA:0 (Tesla T4
Model summary (fused): 168 layers, 3006038 parameters, 0 gradients, 8.1 GFLOP
val: Scanning /content/datasets/klt_detection-1/valid/labels.cache... 19 imag
/usr/lib/python3.10/multiprocessing/popen_fork.py:66: RuntimeWarning: os.fork
    self.pid = os.fork()
          Class      Images   Instances     Box(P)        R      mAP50
            all         19       36       0.912      0.873      0.926
            klt         19       14       0.977      0.929      0.99
            person      19       22       0.846      0.818      0.863
Speed: 5.1ms preprocess, 26.0ms inference, 0.0ms loss, 54.1ms postprocess per
Results saved to runs/detect/val
💡 Learn more at https://docs.ultralytics.com/modes/val
```

▼ Inference with Custom Model

```
%cd {HOME}  
!yolo task=detect mode=predict model={HOME}/runs/detect/train4/weights/best.pt con
```

NOTE: Let's take a look at few results.

```
import glob  
from IPython.display import Image, display  
  
for image_path in glob.glob(f'{HOME}/runs/detect/predict3/*.jpg')[:3]:  
    display(Image(filename=image_path, width=600))  
    print("\n")
```